

10601


Machine Learning

Evaluating classifiers: Model and
feature selection

Model selection issues

- **We have seen some of this before ...**
- **Selecting features (or basis functions)**
 - Linear regression
 - Logistic regression
 - SVMs
- **Selecting parameter value**
 - Prior strength
 - Naïve Bayes, linear and logistic regression
 - Regularization strength
 - Linear and logistic regression
 - Decision trees
 - depth, number of leaves
 - Clustering
 - Number of clusters
- More generally, these are called **Model Selection Problems**

Training and test set error as a function of model complexity



Select points by clicking on the graph or press


Example

Degree of polynomial: ☒ Fit Y to X
☐ Fit X to Y

Calculate

View Polynomial

Reset



Select points by clicking on the graph or press

Example

Degree of polynomial: ☒ Fit Y to X
☐ Fit X to Y

Calculate

View Polynomial

Reset

Simple greedy model selection algorithm

- Pick a dictionary of features
 - e.g., polynomials for linear regression
- Greedy heuristic:
 - Start from empty (or simple) set of features $F_0 = \emptyset$
 - Run learning algorithm for current set of features F_t
 - Obtain h_t
 - Select **next best feature** X_i^*
 - e.g., X_j that results in lowest training error learner when learning with $F_t \cup \{X_j\}$
 - $F_{t+1} \leftarrow F_t \cup \{X_i^*\}$
 - Recurse

Greedy model selection

- Applicable in many settings:
 - Linear regression: Selecting basis functions
 - Naïve Bayes: Selecting (independent) features $P(X_i|Y)$
 - Logistic regression: Selecting features (basis functions)
 - Decision trees: Selecting leaves to expand
- Only a heuristic!
 - But, sometimes you can prove something cool about it

Simple greedy model selection algorithm

- Greedy heuristic:
 - ...
 - Select **next best feature** X_i^*
 - e.g., X_j that results in lowest training error learner when learning with $F_t \cup \{X_j\}$
 - $F_{t+1} \leftarrow F_t \cup \{X_i^*\}$
 - Recurse

When do you stop???

- When training error is low enough?
- When test set error is low enough?



But how can we tell what is the
test error?

Validation set

- In general, for evaluating a classifier we **randomly** split a dataset into two parts:
 - Training data – $\{\mathbf{x}_1, \dots, \mathbf{x}_{N_{\text{train}}}\}$
 - Test data – $\{\mathbf{x}_1, \dots, \mathbf{x}_{N_{\text{test}}}\}$
- Why not use the test data to determine the correct model? Or when to stop?
- **Test data must always remain independent!**
 - Never ever ever ever learn on test data, including for model selection
- Given a dataset, **randomly** split it into three parts:
 - Training data – $\{\mathbf{x}_1, \dots, \mathbf{x}_{N_{\text{train}}}\}$
 - Validation data – $\{\mathbf{x}_1, \dots, \mathbf{x}_{N_{\text{valid}}}\}$
 - Test data – $\{\mathbf{x}_1, \dots, \mathbf{x}_{N_{\text{test}}}\}$
- Use validation data for tuning learning algorithm, e.g., model selection
 - Save test data for very final evaluation

Simple greedy model selection algorithm

- Greedy heuristic:
 - ...
 - Select **next best feature** X_i^*
 - e.g., X_j that results in lowest training error learner when learning with $F_t \cup \{X_j\}$
 - $F_{t+1} \leftarrow F_t \cup \{X_i^*\}$
 - Recurse


When do you stop???

- When training error is low enough?
- When test set error is low enough?
- When validation set error is low enough?

Sometimes, but there is an even better option ...

Validating a learner, not a hypothesis (intuition only, not proof)

- With a validation set, we can estimate the error of 1 hypothesis on 1 dataset
 - e.g. Should I use a polynomial of degree 3 or 4
- However, what we really want is to estimate the error of learner over multiple datasets



(LOO) Leave-one-out cross validation

- Consider a **validation set with 1 example**:
 - D – training data
 - $D \setminus i$ – training data with i th data point moved to validation set
- **Learn classifier $h_{D \setminus i}$ with the $D \setminus i$ dataset**
- **Estimate true error** as:
 - 0 if $h_{D \setminus i}$ classifies i th data point correctly
 - 1 if $h_{D \setminus i}$ is wrong about i th data point
 - Seems really bad estimator, but wait!
- **LOO cross validation**: Average over all data points i :
 - **For each data point you leave out, learn a new classifier $h_{D \setminus i}$**
 - **Estimate error** as:

$$error_{LOO} = \frac{1}{m} \sum_{i=1}^m \mathbb{1} \left(h_{D \setminus i}(\mathbf{x}^i) \neq y^i \right)$$

LOO cross validation is (almost) unbiased estimate of true error!

- When computing **LOOCV error**, we only use *m-1* data points
 - So it's not an estimate of true error of learning with *m* data points
 - Usually pessimistic, though – learning with less data typically gives worse answer
- **LOO is almost unbiased!**
 - Let $error_{true,m-1}$ be true error of learner when you only get *m-1* data points
 - LOO is unbiased estimate of $error_{true,m-1}$:

$$E_{\mathcal{D}}[error_{LOO}] = error_{true,m-1}$$

- **Great news!**
 - Use **LOO error** for model selection!!!


Simple greedy model selection algorithm

- Greedy heuristic:
 - ...
 - Select **next best feature** X_i^*
 - e.g., X_j that results in lowest training error learner when learning with $F_t \cup \{X_j\}$
 - $F_{t+1} \leftarrow F_t \cup \{X_i^*\}$
 - Recurse

When do you stop???

- When training error is low enough?
- When test set error is low enough?
- When validation set error is low enough?
- **STOP WHEN error_{LOO} IS LOW!!!**


LOO cross validation error



Select points by clicking on the graph or press

Example

Degree of polynomial: ☒ Fit Y to X
☐ Fit X to Y



Select points by clicking on the graph or press

Example

Degree of polynomial: ☒ Fit Y to X
☐ Fit X to Y

Computational cost of LOO

- Suppose you have 100,000 data points
- You implemented a great version of your learning algorithm
 - Learns in only 1 second
- Computing LOO will take about 1 day!!!
 - If you have to do for each choice of basis functions, it will take forever!

Solution: Use k -fold cross validation

- Randomly **divide training data into k equal parts**

- D_1, \dots, D_k

- For each i

- **Learn classifier $h_{D \setminus D_i}$ using data point not in D_i**

- **Estimate error of $h_{D \setminus D_i}$ on validation set D_i :**

$$error_{\mathcal{D}_i} = \frac{k}{m} \sum_{(\mathbf{x}^j, y^j) \in \mathcal{D}_i} \mathbb{1}(h_{D \setminus \mathcal{D}_i}(\mathbf{x}^j) \neq y^j)$$

- **k -fold cross validation error is average** over data splits:

$$error_{k-fold} = \frac{1}{k} \sum_{i=1}^k error_{\mathcal{D}_i}$$

- k -fold cross validation properties:

- **Much faster to compute** than LOO

- **More (pessimistically) biased** – using much less data, only $m(k-1)/k$

Regularization


- Model selection 1 (using CV): **Greedy**
 - Pick subset of features that have yield low LOO error
- Model selection 2: **Regularization**
 - Include **all possible features!**
 - **Penalize “complicated” hypothesis**

Regularization in linear regression

- Overfitting usually leads to very large parameter choices, e.g.:

$$-2.2 + 3.1 X - 0.30 X^2$$

$$-1.1 + 4,700,910.7 X - 8,585,638.4 X^2 + \dots$$



- Regularized least-squares (a.k.a. ridge regression):

$$w^* = \arg \min_w \sum_j (w^T x_j - y_j)^2 + \lambda \sum_i w_i^2$$

Other regularization examples

- **Logistic regression** regularization
 - Maximize data likelihood minus **penalty for large parameters**

$$\arg \max_{\mathbf{w}} \sum_j \ln P(y^j | \mathbf{x}^j, \mathbf{w}) - \lambda \sum_i w_i^2$$

- **Biases towards small parameter values**
-
- **Naïve Bayes** regularization
 - **Prior** over likelihood of features
 - **Biases away from zero probability** outcomes
-
- **Decision tree** regularization
 - Many possibilities, e.g., **Chi-Square test**
 - **Biases towards smaller trees**
-
- **Sparsity**: find good solution with few basis functions, e.g.:
 - **Simple greedy model selection** from earlier in the lecture
 - **L1 regularization, e.g.:**

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_j (\mathbf{w}^T \mathbf{x}_j - y_j)^2 + \lambda \sum_i |w_i|$$

Regularization and Bayesian learning

$$p(\mathbf{w} \mid Y, \mathbf{X}) \propto P(Y \mid \mathbf{X}, \mathbf{w})p(\mathbf{w})$$

- For example, if we assume a zero mean, Gaussian prior for w in a logistic regression classification we would end up with an L2 regularization
 - Why?
 - Board ...
 - What is λ ?
- Similar interpretation for other learning approaches:
 - **Linear regression**: Also zero mean, Gaussian prior for \mathbf{w}
 - **Naïve Bayes**: Directly defined as prior over parameters

How do we pick magic parameter λ ?

Cross Validation!!!

Occam's Razor



- William of Ockham (1285-1349) *Principle of Parsimony*:
 - “One should not increase, beyond what is necessary, the number of entities required to explain anything.”
- Regularization penalizes for “*complex explanations*”
- Alternatively (but pretty much the same), use *Minimum Description Length (MDL) Principle*:
 - minimize $length(\text{misclassifications}) + length(\text{hypothesis})$
- $length(\text{misclassifications})$ – e.g., #wrong training examples
- $length(\text{hypothesis})$ – e.g., size of decision tree

Minimum Description Length Principle

- MDL prefers small hypothesis that fit data well:

$$h_{MDL} = \arg \min_h L_{C_1}(\mathcal{D} | h) + L_{C_2}(h)$$

- $L_{C_1}(D|h)$ – description length of data under code C_1 given h
 - Only need to describe points that h doesn't explain (classify correctly)
- $L_{C_2}(h)$ – description length of hypothesis h
- Decision tree example
 - $L_{C_1}(D|h)$ – #bits required to describe data given h
 - If all points correctly classified, $L_{C_1}(D|h) = 0$
 - $L_{C_2}(h)$ – #bits necessary to encode tree
 - Trade off quality of classification with tree size


What you need to know about Model Selection, Regularization and Cross Validation

- Cross validation
 - (Mostly) Unbiased estimate of true error
 - LOOCV is great, but hard to compute
 - k -fold much more practical
 - Use for selecting parameter values!
- Model selection
 - Search for a model with low cross validation error
- Regularization
 - Penalizes for complex models
 - Select parameter with cross validation
 - Really a Bayesian approach
- Minimum description length
 - Information theoretic interpretation of regularization

Feature selection

- Choose an optimal subset from the set of all N features
 - Only use a subset of a possible words in a dictionary
 - Only use a subset of genes
- Why?
- Can we do model selection to solve this? – 2^n models

eg. Microarray data



Courtesy : Paterson
Institute

Two approaches: 1. Filter

- Independent of classifier used
- Rank features using some criteria based on their relevance to the classification task
- For example, mutual information:

$$I(X; Y) = \sum_{y \in Y} \sum_{x \in X} p(x, y) \log \left(\frac{p(x, y)}{p_1(x) p_2(y)} \right),$$

- Choose a subset based on the sorted scores for the criteria used

2. Wrapper

- Classifier specific
- Greedy (large search space)
- Initialize $F = \text{null set}$
 - At each step, using cross validation or an information theoretic criteria, choose a feature to add to the subset [training should be done with only features in $F + \text{new feature}$]
 - Add the chosen feature to the subset
- Repeat until no improvement to CV accuracy